


XGBoost Classifier Architecture

Gradient Boosted Trees (1,000 rounds, depth 10)



Each tree fits the negative gradient (residual) of the previous ensemble

$$F_{t+1}(x) = F_t(x) + \eta \cdot h_{t+1}(x)$$

$$\text{Objective: } \mathcal{L} = \sum_i L(y_i, \hat{y}_i) + \sum_k \lambda \sum_j w_j^2$$

Hyperparameters (shared by FA and WD classifiers)

